

HTTP Server Code Review Exam

Interviewer Name

April 6, 2025

1 Architecture and Design (30 points)

1. What threading model does this HTTP server implement? Explain the roles of:
 - The listener thread
 - The worker threads
2. Why does the server use epoll for I/O multiplexing instead of:
 - `select()`
 - `poll()`

What are the advantages of epoll in this context?

3. The server uses a round-robin approach to distribute connections to worker threads.
 - What are the pros and cons of this approach?
 - What alternative load balancing strategies could be used?
4. Explain the purpose of the `EventData` structure and how it's used in the event handling cycle.

2 Implementation Details (40 points)

5. The `HandleEpollEvent` method handles both `EPOLLIN` and `EPOLLOUT` events.
 - Describe the complete lifecycle of a single HTTP request through this method
 - Why does it need to manage memory allocation/deallocation of `EventData` structures?
6. The server uses non-blocking sockets with edge-triggered epoll.
 - What modifications would be needed to switch to level-triggered mode?
 - What are the implications of this change?
7. Examine the error handling in the `ProcessEvents` method:
 - What conditions cause a socket to be closed?
 - Why does it check for both `EPOLLHUP` and `EPOLLERR`?
8. The `control_epoll_event` helper function handles three epoll operations:
 - What are the race conditions that could occur when modifying epoll sets?
 - How does the current implementation prevent them?

3 Performance and Optimization (20 points)

9. The server implements an exponential backoff mechanism using `sleep_times_`.
 - When is this used and why?
 - How could this be improved?
10. The current implementation copies HTTP response data twice:
 - Where do these copies occur?
 - How could you reduce this to a single copy?

11. The server uses a fixed-size thread pool.
- What are the limitations of this approach?
 - How could you make it dynamically scalable?

4 Advanced Questions (10 points)

12. How would you modify this server to support HTTPS/SSL connections?
- What components would need to change?
 - How would this affect performance?
13. The current URI-based routing is simple. How would you implement:
- Path parameters (e.g., `/users/:id`)
 - Middleware support
14. How could you add connection timeout handling?
- For idle connections
 - For slow request/response cycles

Scoring Guide

- 90-100: Expert-level understanding
- 70-89: Strong working knowledge
- 50-69: Basic understanding with some gaps
- Below 50: Needs significant study