

HTTP Server Code Review - Answers

Candidate Name

April 6, 2025

1 Architecture and Design Answers

1. Threading Model:

- *Listener thread*: Accepts new connections and distributes them to worker threads using round-robin. Runs in a loop while `running_` is true.
- *Worker threads*: Each has its own epoll instance to handle I/O events for assigned connections. Process both read (EPOLLIN) and write (EPOLLOUT) events.

2. epoll Advantages:

- `select()/poll()` scale $O(n)$ with connections, while epoll scales $O(1)$
- epoll uses edge-triggered mode which avoids repeated notifications
- Better performance with large numbers of connections (10k+)
- Kernel maintains ready list instead of scanning all fds

3. Round-Robin Distribution:

- Pros: Simple, no synchronization needed between workers
- Cons: Can lead to imbalance if connections have uneven workloads
- Alternatives: Least connections, weighted round-robin, or work stealing queues

4. EventData Purpose:

- Carries buffer state (cursor/length for partial reads/writes)
- Preserves context between epoll events (request/response cycle)
- Contains file descriptor and buffer for I/O operations
- Lifetime managed between EPOLLIN → EPOLLOUT transitions

2 Implementation Details Answers

5. `HandleEpollEvent` Lifecycle:

- EPOLLIN: Reads request → Parses → Generates response → Switches to EPOLLOUT
- EPOLLOUT: Writes response → Returns to EPOLLIN or closes connection
- Memory Management: Allocates new `EventData` for each state transition to prevent use-after-free

6. Level-Triggered Changes:

- Remove EPOLLET flag from event registration
- Must handle repeated notifications for same data
- Simpler but less efficient (more syscalls)

7. Error Handling:

- Close Conditions: EOF (`recv=0`), errors (`EPOLLHUP/EPOLLERR`), send failures
- `EPOLLHUP`=hangup (remote closed), `EPOLLERR`=error (local problem)
- Both indicate unrecoverable connection states

8. `control_epoll_event` Race Conditions:

- Potential FD reuse between delete/add
- Current solution: Serialized through single-threaded event processing
- Better: Could use thread-safe FD tracking

3 Performance and Optimization Answers

9. Exponential Backoff:

- Used when `accept()` or `epoll_wait()` return no events
- Prevents CPU spinning during idle periods
- Improvement: Adaptive sleeping based on load metrics

10. Copy Elimination:

- Current copies:
 - (a) `recv()` → request buffer
 - (b) response string → send buffer
- Optimization: Use scatter/gather I/O (`writv`) or response buffers directly

11. Thread Pool Limitations:

- Fixed size can't adapt to load spikes
- Dynamic scaling: Monitor queue depth, add/remove workers
- Challenge: Epoll FDs aren't thread-safe to share

4 Advanced Questions Answers

12. HTTPS Support:

- Changes Needed:
 - `SSL_accept/SSL_read/SSL_write` wrappers
 - TLS handshake state machine
 - Certificate management
- Performance Impact: 10x slower due to crypto ops

13. Routing Enhancements:

- Path Parameters: Use regex matching with capture groups
- Middleware: Chain of responsibility pattern with handler stack

14. **Timeout Handling:**

- **Idle:** Timerfd + epoll with connection timestamp checks
- **Slow requests:** Per-operation deadlines with SO_RCVTIMEO/SO_SNDTIMEO

Code Improvements

- **Memory:** Use memory pools for EventData allocations
- **Logging:** Add connection lifecycle tracing
- **Stats:** Track request rates/latencies per worker
- **HTTP/2:** Frame-based protocol support